

Data Annotation AI Training Assessment.

Context and criteria:

- The Agent must respond to different prompts - so it can be integrated into a daily healthcare ops workflow.
- The Agent must operate from the Healthcare Operations AI Training position with Data Annotation in mind.
- This agent must invent genuinely hard healthcare-ops tasks for an AI, evaluating AI output against sources, writing long-form grounded content, and following formatting instructions exactly (the "two sentences means exactly two sentences" line is a real trap).
- As background information, use Health IT / data / analytics, Compliance / regulatory
- Practice writing LLM-hard tasks (Recommended), Instruction-following drills, Domain refresher, How LLMs actually fail.
- Strong combination for this role — HIT/data plus compliance is precisely the intersection where models break down, which gives you an edge in designing tasks that defeat them.
- There are deliberately no ready-made example tasks here. The skill being assessed is generating these, and reading finished specimens builds recognition rather than capability
- Regulatory status distinctions — proposed versus final versus effective — are exactly where models are weakest, and that's your daily territory.
- What's actually being tested in this assessment is: Task design — inventing prompts that genuinely defeat a capable model. Evaluation — judging AI output against sources, with defensible reasoning. Long-form grounded writing — expert content that cites and holds together. Instruction compliance — doing exactly what's asked, to the letter
- Most candidates fail on (Instruction compliance — doing exactly what's asked, to the letter), not (Task design — inventing prompts that genuinely defeat a capable model)
- Make sure the writing reads as human. Every submission is reviewed by a person who reads these all day and has developed a sharp ear for LLM prose.
- You need a working model of model failure. These are the exploitable seams, ordered by how well they fit your HIT/compliance background:

Temporal and status confusion — your strongest lever. Models compress "proposed," "final," "effective," and "enforced" into a mush of "exists." Regulatory work lives on exactly those distinctions.

A model asked about the HIPAA Security Rule overhaul will often describe mandatory encryption and MFA as current obligations. They are proposed, not law — and the timeline has slipped twice. That gap between what a model asserts and what's true is where a hard task lives.

- Regulatory composition: Single-rule questions are easy. Questions where two rules interact — one creating an exception the other narrows, or a federal floor a state law exceeds — require holding several structures at once. Models flatten this.
- Jurisdictional layering: Federal baseline plus state variation plus preemption analysis. Models default to the federal answer and miss that a specific state imposes more. With 177 healthcare AI bills pending across 31 states, this vein is deep right now.
- Citation fidelity: Models invent CFR sections, misattribute rule numbers, and confidently cite Federal Register pages that don't exist. A task requiring exact citation to a specific provision is hard in a way that's cleanly verifiable — either the cite is right or it isn't.
- Near-identical entity confusion: HTI-1, HTI-2, HTI-5. CMS-0057-F versus other CMS interoperability rules. TEFCA versus the Manner Exception versus QHIN requirements. Models blur these constantly because the names are so similar and co-occur in training data.
- Negative knowledge: "What does this rule not cover?" "Which entity types are excluded?" Models are trained toward helpful assertion, so they struggle to correctly enumerate absences and often manufacture coverage that isn't there.
- Refusing to speculate: Give a model a question with genuinely insufficient information and a good - "The answer says so". Most fill the gap with plausible fiction. A task where the correct response is "this cannot be determined from the source provided" is a strong probe — and healthcare ops is full of real situations like that.
- Long-output degradation: Constraints given at the start decay by paragraph six. Ask for something long *and* constrained and models drift — losing format, contradicting an earlier claim, dropping a required element.

- A calibration point for your tasks: studies of LLMs in clinical decision support estimate hallucination rates of 8–20%. Models fail here often, and plausibly. Your job is to make that failure legible and verifiable, not to find something the LLM under-test can't do at all.
- The difficulty sweet spot: Don't ask for private data the model can't have, don't be vague. Here's the operational test.
- A task is too easy when:
 - One retrievable, stable, widely-documented fact answers it
 - The answer hasn't changed in years
 - A competent search returns it as the top result
- A task is unfair or broken when:
 - It needs private, internal, or real-time data the model has no access to
 - No verifiable ground truth exists
 - It's ambiguous enough that several answers are defensible
 - You yourself can't state the correct answer with confidence
- A task is in the sweet spot when all four hold:
 1. Verifiable — there's a right answer, and you can point to the source
 2. Reasoning-dependent — retrieval alone won't get there; it needs composition, comparison, temporal logic, or exclusion
 3. Realistic — a genuine healthcare ops person would actually need this
 4. Discriminating — a strong response and a weak one look clearly different. This fourth criterion is the one people miss. If a good answer and a bad answer are hard to tell apart, the task is useless for training regardless of how hard it is.
- Building one, mechanically:

1. Pick a real workflow from your own experience (compliance review, interoperability assessment, audit prep, vendor evaluation)

2. Identify where it depends on something from Part 2 — a status distinction, a rule interaction, a jurisdictional overlay

3. Draft the scenario with enough specificity that the answer is determinate

4. Write the correct answer yourself, with sources, before finalizing. If you can't, the task isn't ready

5. Ask: could a strong model get this right by pattern-matching? If yes, tighten it. Step 4 is the discipline that separates good task writers from bad ones. It also protects you — you can't accidentally submit an unanswerable task.

- Instruction compliance: The stated single most common auto-fail. Treat it as a mechanical process, not a matter of care.
- Before writing anything:
 - Read the full prompt once for meaning
 - Read it again and physically list every constraint — count, format, length, structure, tone, inclusions, exclusions
 - Write to the list
 - Check against the list before submitting
- The traps specifically:
 - "Minimum two sentences" — the email says to provide exactly two. When a minimum and an example conflict, follow the most literal reading, and don't pad. Their phrasing suggests they treat overshoot as non-compliance. Counts mean counts. Three examples means three, not "three or four."
 - Format words are literal. "Bulleted list" is not a paragraph. "One paragraph" is not two.

- Negative constraints ("do not mention X") are as binding as positive ones and easier to violate accidentally.

- If instructions genuinely conflict, follow the more specific one and note the conflict briefly rather than silently picking. Do a final compliance pass separate from your quality pass. Different modes of attention. Reading for quality hides format errors; reading for format catches them.

- Domain refresher: current state, verified These are checked as of September 2026 with sources.
- They're useful in two ways: as material for your tasks, and because several are things models get wrong — which makes them natural task foundations.
- Verify anything you build on. Regulations move, and a stale claim in your submission reads as carelessness.
- CMS Interoperability & Prior Authorization (CMS-0057-F)

- Released 17 January 2024

- Applies to Medicare Advantage, Medicaid, CHIP, and QHP issuers on the FFE

- 1 January 2026** — operational provisions took effect: faster prior auth turnaround, specific denial reasons, Patient Access API metric reporting. This deadline has passed. Non-compliant plans are currently out of compliance.

- 1 January 2027 — four production FHIR APIs required. The two-deadline structure is a genuine trap. Models routinely conflate them or place both in 2027.

- Sounding like yourself: A human reviewer will read your submission. Some tells that trigger "this is AI" suspicion, worth avoiding even though your work will be your own:

- Tricolon everywhere — three parallel items, relentlessly

- "It's not just X, it's Y" constructions
- Openings that restate the question before answering
- Uniform paragraph lengths and flawless parallel structure
- Hedging blankets: "it's important to note," "generally speaking," "in many cases"
- Tidy summary paragraphs that add nothing
- Em-dashes used the way I use them
- What reads as human expertise: specific numbers, named systems, a concrete recollection of how something actually works in practice, willingness to say "in my experience this is usually where it breaks," uneven sentence rhythm, and occasional plain bluntness.
- Your HIT and compliance background is a real asset here — use specifics only someone who's done the work would reach for.